


# EECE7205: Fundamentals of Computer Engineering



## Elementary Data Structures



# Data Structures

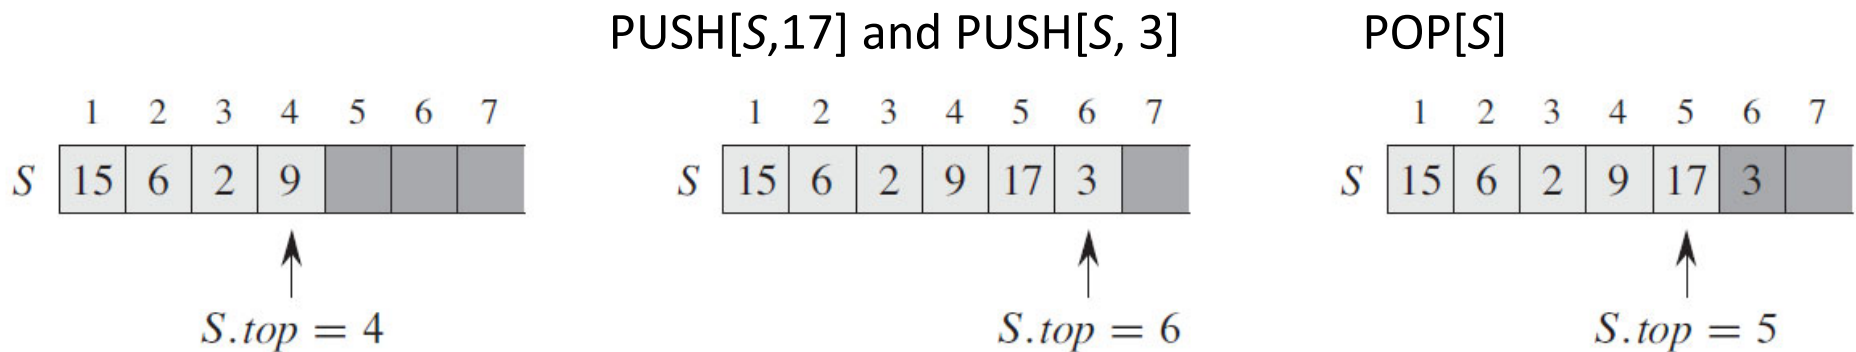
---

- A ***data structure*** is a way to store and organize data in order to facilitate access and modifications.
- No single data structure works well for all purposes, and so it is important to know the strengths and limitations of several of them.
- We will cover different data structures in this course.



# Stacks

- Stacks are dynamic sets in which the element removed from the set is the one most recently inserted.
- The stack implements a ***last-in, first-out***, or ***LIFO***, policy.
- The INSERT operation on a stack is often called PUSH, and the DELETE operation is often called POP.
- The following figure shows that we can implement a stack of at most  $n$  elements with an array  $S[1..n]$ . The array has an attribute  $S.top$  that indexes the most recently inserted element.
- When  $S.top = 0$ , the stack contains no elements and is ***empty***.





# Stack Operations

- Each of the following three stack operations takes  $O(1)$  time.
- Do we need to check for the array capacity in PUSH?

STACK-EMPTY( $S$ )

```

1  if  $S.top == 0$ 
2      return TRUE
3  else return FALSE

```

PUSH( $S, x$ )

```

1   $S.top = S.top + 1$ 
2   $S[S.top] = x$ 


```

POP( $S$ )

```

1  if STACK-EMPTY( $S$ )
2      error "underflow"
3  else  $S.top = S.top - 1$ 
4      return  $S[S.top + 1]$ 

```






# Queues


- Queues are dynamic sets in which the element removed from the set is the one that has been in the set for the longest time.
- The queue implements a *first-in, first-out*, or **FIFO**, policy.
- We call the INSERT operation on a queue ENQUEUE, and we call the DELETE operation DEQUEUE.
- The FIFO property of a queue causes it to operate like a line of customers waiting to pay a cashier.
- The queue has a *head* and a *tail*.
- When an element is enqueued, it takes its place at the tail of the queue.
- The element dequeued is always the one at the head of the queue.




# Queue Example



ENQUEUE( $Q$ , 17),  
ENQUEUE( $Q$ , 3), and  
ENQUEUE( $Q$ , 5)



DEQUEUE( $Q$ )





# Queue Operations

- Initially, we have  $Q.head = Q.tail = 1$

## ENQUEUE( $Q, x$ )

```

1   $Q[Q.tail] = x$ 
2  if  $Q.tail == Q.length$ 
3       $Q.tail = 1$ 
4  else  $Q.tail = Q.tail + 1$ 


```

## DEQUEUE( $Q$ )

```

1   $x = Q[Q.head]$ 
2  if  $Q.head == Q.length$ 
3       $Q.head = 1$ 
4  else  $Q.head = Q.head + 1$ 
5  return  $x$ 

```




- When  $Q.head == Q.tail$ , is the queue empty or full?



# Linked Lists

- A **linked list** is a data structure in which the objects are arranged in a linear order.
  - Unlike an array, however, in which the linear order is determined by the array indices, the order in a linked list is determined by pointers in each object.
- As shown, each element of a **doubly linked list**  $L$  is an object with an attribute *key* and two other pointer attributes: *next* and *prev*.




- If a list is **singly linked**, we omit the *prev* pointer in each element.





# Doubly Linked List

- In a **doubly linked list**  $L$ :
  - If  $x.prev = \text{NIL}$ , the element  $x$  has no predecessor and is therefore the first element, or **head**, of the list.
  - If  $x.next = \text{NIL}$ , the element  $x$  has no successor and is therefore the last element, or **tail**, of the list.
  - An attribute  $L.head$  points to the first element of the list.
  - If  $L.head = \text{NIL}$ , the list is empty.






# Doubly Linked List Node

- Each node in a **doubly linked list** contains not only the values of the elements we need to store in the list, like a Queue data structure, but also the prev and next pointers.
- A node in a list of integers can be defined as:

```
struct Node
{
    Node* prev;
    int    item;
    Node* next; };

```

- A node is created dynamically as:  
`Node* x = new Node;`
- Where **x** is used to insert the new node in the list.





# Searching a Linked List


- The procedure  $\text{LIST-SEARCH}(L, k)$  finds the first element with key  $k$  in list  $L$  by a simple linear search, returning a pointer to this element.
- If no object with key  $k$  appears in the list, then the procedure returns NIL.

$\text{LIST-SEARCH}(L, k)$

```

1   $x = L.\text{head}$ 
2  while  $x \neq \text{NIL}$  and  $x.\text{key} \neq k$ 
3       $x = x.\text{next}$ 
4  return  $x$ 

```





# Inserting Into a Linked List


- Given an element  $x$  whose *key* attribute has already been set, the LIST-INSERT procedure inserts  $x$  onto the front of the linked list.

LIST-INSERT( $L, x$ )

```

1   $x.next = L.head$ 
2  if  $L.head \neq \text{NIL}$ 
3       $L.head.prev = x$ 
4   $L.head = x$ 
5   $x.prev = \text{NIL}$ 

```



Following the execution of LIST-INSERT( $L, x$ ), where  $x.key = 25$



# Deleting From a Linked List

- The procedure LIST-DELETE removes an element  $x$  from a linked list  $L$ . It must be given a pointer to  $x$ , and it then delete  $x$  out of the list by updating pointers.

**LIST-DELETE** ( $L, x$ )

```

1  if  $x.prev \neq \text{NIL}$ 
2       $x.prev.next = x.next$ 
3  else  $L.head = x.next$ 
4  if  $x.next \neq \text{NIL}$ 
5       $x.next.prev = x.prev$ 
  
```



Following the execution of **LIST-DELETE**( $L, x$ ), where  $x.key = 4$



# Linked List Sentinels


- The code for LIST-DELETE would be simpler if we could ignore the boundary conditions at the head and tail of the list:

LIST-DELETE' ( $L, x$ )

1  $x.\text{prev}.\text{next} = x.\text{next}$

2  $x.\text{next}.\text{prev} = x.\text{prev}$

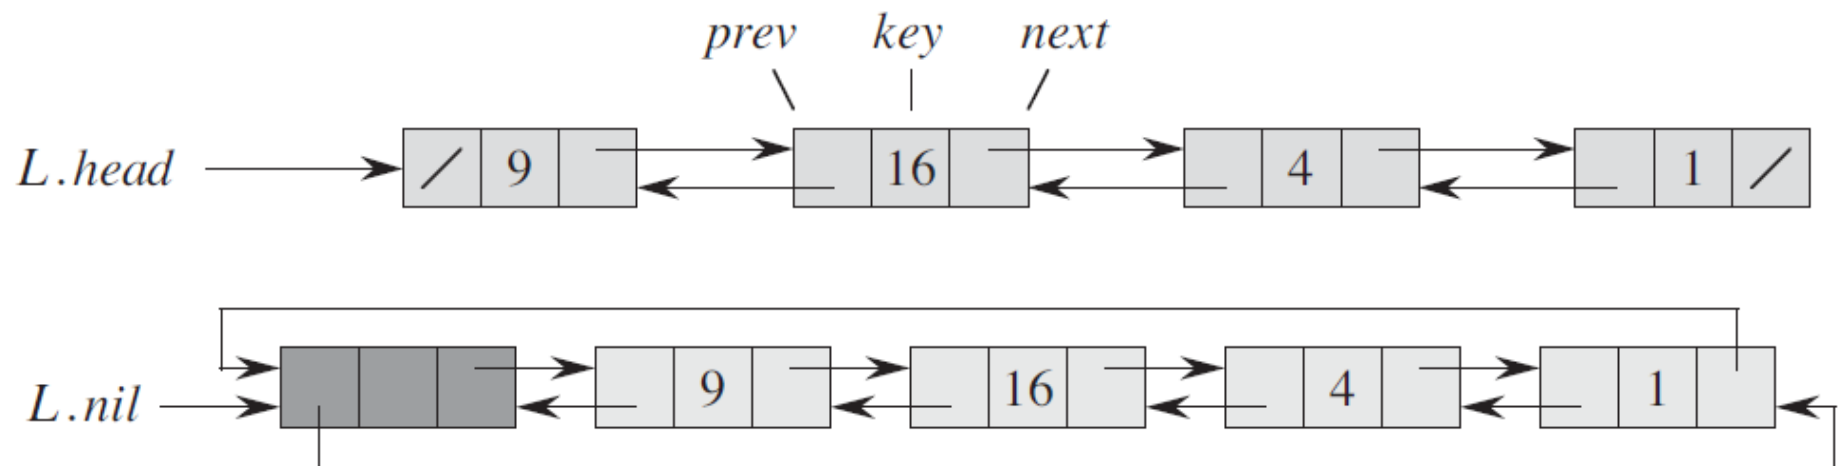
- One way to achieve that is to use a circular, doubly linked list with a sentinel.
  - A **sentinel** is a dummy object that allows us to simplify boundary conditions. The **sentinel** is  $L.\text{nil}$  in the following example:





# A Circular, Doubly Linked List With a Sentinel

- In a circular, doubly linked list with a sentinel, the sentinel *L.nil* lies between the head and tail.
- The attribute *L.nil.next* points to the head of the list, and *L.nil.prev* points to the tail.
- Similarly, both the *next* attribute of the tail and the *prev* attribute of the head point to *L.nil*.





# Searching and Inserting

- The gain from using sentinels within loops is usually a matter of clarity of code rather than speed.
- The linked list code, for example, becomes simpler when we use sentinels, but we save only  $O(1)$  time in the LIST-INSERT' and LIST-DELETE' procedures.
- We should use sentinels judiciously. When there are many small lists, the extra storage used by their sentinels can represent significant wasted memory.

## LIST-SEARCH'( $L, k$ )

```

1   $x = L.nil.next$ 
2  while  $x \neq L.nil$  and  $x.key \neq k$ 
3       $x = x.next$ 
4  return  $x$ 
  
```

## LIST-INSERT'( $L, x$ )

```

1   $x.next = L.nil.next$ 
2   $L.nil.next.prev = x$ 
3   $L.nil.next = x$ 
4   $x.prev = L.nil$ 
  
```






# Binary Trees


- The methods we have studied for representing lists can be extended to represent rooted trees data structure.
- In the following representation of a binary tree  $T$ , each node  $x$  has the attributes  $x.p$  (top),  $x.left$  (lower left), and  $x.right$  (lower right). The key attributes are not shown.

- The root of the entire tree  $T$  is pointed to by the attribute  $T.root$ .
- If  $T.root = \text{NIL}$ , then the tree is empty.
- If  $x.p = \text{NIL}$ , then  $x$  is the root.



# Direct-Address Tables

- Direct addressing is a simple technique that works well when the universe  $U$  of keys is reasonably small assuming that no two elements have the same key.





# Direct-Address Table Operations

- Each of the following operations takes only  $O(1)$  time.

DIRECT-ADDRESS-SEARCH( $T, k$ )

1 **return**  $T[k]$

DIRECT-ADDRESS-INSERT( $T, x$ )

1  $T[x.key] = x$

DIRECT-ADDRESS-DELETE( $T, x$ )

1  $T[x.key] = \text{NIL}$




# Hash Tables

- The downside of direct addressing is obvious: if the universe  $U$  is large, storing a table  $T$  of size  $|U|$  may be impractical, or even impossible, given the memory available on a typical computer.
- Furthermore, the set  $K$  of keys *actually stored* may be so small relative to  $U$  that most of the space allocated for  $T$  would be wasted.
- When the number of keys actually stored is small relative to the total number of possible keys, **hash tables** become an effective alternative to directly addressing an array.
- A hash table is an effective data structure for implementing ***dictionaries*** operations INSERT, SEARCH, and DELETE.
- Although searching for an element in a hash table can take as long as searching for an element in a linked list –  $\Theta(n)$  time in the worst case - in practice, hashing performs extremely well.



# Hash Functions

- With direct addressing, an element with key  $k$  is stored in slot  $k$ .
- With hashing, this element is stored in slot  $h(k)$ ; that is, we use a **hash function**  $h$  to compute the slot from the key  $k$ .
- A hash function  $h$  must be **deterministic** in that a given input  $k$  should always produce the same output  $h(k)$ .
- In the shown example,  $h$  maps the universe  $U$  of keys into the slots of a **hash table**  $T[0 .. m - 1]$  where the size  $m$  of the hash table is typically much less than  $|U|$





# Hash Collision


---

- Two keys may hash to the same slot. We call this situation a ***collision***.
- The ideal solution would be to avoid collisions altogether or at least minimizing their number.
- We might try to achieve this goal by choosing a suitable hash function  $h$ .
- Because  $|U| > m$ , however, there must be at least two keys that have the same hash value; avoiding collisions altogether is therefore impossible.



# Collision Resolution by Chaining

- In **chaining**, we place all the elements that hash to the same slot into the same linked list, as shown.





# Hash Table Operations

- The dictionary operations on a hash table  $T$  are easy to implement when collisions are resolved by chaining:

CHAINED-HASH-INSERT( $T, x$ )

1 insert  $x$  at the head of list  $T[h(x.key)]$

CHAINED-HASH-SEARCH( $T, k$ )

1 search for an element with key  $k$  in list  $T[h(k)]$

CHAINED-HASH-DELETE( $T, x$ )

1 delete  $x$  from the list  $T[h(x.key)]$





# Hash Functions (Cont'd)

- A good hash function satisfies (approximately) the assumption of simple uniform hashing: each key is equally likely to hash to any of the  $m$  slots, independently of where any other key has hashed to.
- Most hash functions assume that the universe of keys is the set  $\mathbb{N} = \{0, 1, 2, \dots\}$  of natural numbers.
- Thus, if the keys are not natural numbers, we find a way to interpret them as natural numbers.



# Division Hash Function

- In the ***division method*** for creating hash functions, we map a key  $k$  into one of  $m$  slots by taking the remainder of  $k$  divided by  $m$ . That is, the hash function is  $h(k) = k \bmod m$ .
- For example, if the hash table has size  $m = 12$  and the key is  $k = 100$ , then  $h(k) = 4$ .
- Example of bad choice of  $m$  is  $2^p$ , as  $h(k) = p$  lowest-order bits of  $k$ .
- We are better off designing the hash function to depend on all the bits of the key.
- A prime not too close to an exact power of 2 is often a good choice for  $m$ . Example:
  - Suppose we wish to allocate a hash table, with collisions resolved by chaining, to hold elements with keys range from 0 to 2000. We don't mind examining an average of 3 elements in an unsuccessful search, and so we allocate a hash table of size  $m = 701$  because it is a prime near  $2000/3$  but not near any power of 2.



# Resolving Collisions by Open Addressing

- In ***open addressing***, all elements occupy the hash table itself. That is, each table entry contains either an element of the dynamic set or NIL.
- When searching for an element, we systematically examine table slots until either we find the desired element, or we have ascertained that the element is not in the table.
- Unlike chaining, no lists and no elements are stored outside the table.
- Thus, in open addressing, the hash table can “**fill up**” so that no further insertions can be made.
- the advantage of open addressing is that it **avoids pointers** altogether.
- Instead of following pointers, we *compute* the sequence of slots to be examined.



# Open Addressing (Cont'd)

- To perform insertion using open addressing, we successively examine, or **probe**, the hash table until we find an empty slot in which to put the key.
- To determine which slots to probe, we extend the hash function to include the probe number (starting from 0) as a second input.
- For every key  $k$ , The probe sequence  
 $\langle h(k,0), h(k,1), \dots, h(k,m-1) \rangle$   
should be a permutation of  $\{0, 1, \dots, m-1\}$ .
- So that every hash-table position is eventually considered as a slot for a new key as the table fills up.



# Open Addressing - Insert

- The following HASH-INSERT procedure takes as input a hash table  $T$  and a key  $k$ . It either returns the slot number where it stores key  $k$  or flags an error because the hash table is already full.

**HASH-INSERT**( $T, k$ )

```
1   $i = 0$ 
2  repeat
3       $j = h(k, i)$ 
4      if  $T[j] == \text{NIL}$ 
5           $T[j] = k$ 
6          return  $j$ 
7      else  $i = i + 1$ 
8  until  $i == m$ 
9  error "hash table overflow"
```



# Open Addressing - Search

- The algorithm for searching for key  $k$  probes the same sequence of slots that the insertion algorithm examined when key  $k$  was inserted.
- Therefore, the search can terminate (unsuccessfully) when it finds an empty slot.  
(This argument assumes that keys **were not deleted** from the hash table.)

## **HASH-SEARCH**( $T, k$ )

```
1   $i = 0$ 
2  repeat
3       $j = h(k, i)$ 
4      if  $T[j] == k$ 
5          return  $j$ 
6       $i = i + 1$ 
7  until  $T[j] == \text{NIL}$  or  $i == m$ 
8  return NIL
```



# Open Addressing - Delete

- Deletion from an open-address hash table is difficult. When we delete a key from slot  $i$ , we cannot simply mark that slot as empty by storing NIL in it. If we did, we might be unable to retrieve any key  $k$  during its insertion we had probed slot  $i$  and found it occupied.
- We can solve this problem by marking the slot, storing in it the **special value DELETED** instead of **NIL**.
- We would then modify the procedure HASH-INSERT to treat such a slot as if it were empty so that we can insert a new key there.
- We do not need to modify HASH-SEARCH, since it will pass over DELETED values while searching.



# Linear Probing

- Given an ordinary hash function  $h'(k)$ , linear probing uses the hash function:

$$h(k,i) = (h'(k) + i) \bmod m \text{ for } i = 0, 1, \dots, m-1.$$

- Given key  $k$ , we first probe  $T[h'(k)]$ , i.e., the slot given by the auxiliary hash function.
- We next probe slot  $T[h'(k) + 1]$ , and so on up to slot  $T[m - 1]$  and then we wrap around to slots  $T[0]$ ,  $T[1]$ , ... until we finally probe slot  $T[h'(k) - 1]$ . Because the initial probe determines the entire probe sequence, there are only  $m$  distinct probe sequences.

