


# EECE7205: Fundamentals of Computer Engineering




## Binary and Balanced Search Trees




# Binary Tree Types


- In a **binary** tree, every node has 0, 1, or 2 children.
  - Nodes with 0 children all called **leaves**.
- In a **full** binary tree, every node has either 0 or 2 children.
- A **complete** binary tree is completely filled on all levels except the lowest level, which is filled from the left to right.
- A **perfect** binary tree is a full binary tree with all **leaves** at the same depth.



**full**



**complete**





**perfect**



# Binary Search Trees (BST)

- A binary search tree is organized, as the name suggests, in a binary tree.
- For any node  $x$ , the keys in the left subtree of  $x$  are at most  $x.key$ , and the keys in the right subtree of  $x$  are at least  $x.key$ .
- Different binary search trees can represent the same set of values (as shown where the top representation is more efficient than the bottom one)





# BST Property

---

Let  $x$  be a node in a binary search tree.


If  $y$  is a node in the left subtree of  $x$ ,  
then  $y.key \leq x.key$ .

If  $y$  is a node in the right subtree of  $x$ ,  
then  $y.key \geq x.key$



# BST Implementation

- We can represent such a tree by a **linked data structure** in which each node is an object.
  - In addition to a *key* and satellite data, each node contains attributes *left*, *right*, and *p* that point to the nodes corresponding to its left child, its right child, and its parent, respectively.
  - If a child or the parent is missing, the appropriate attribute contains the value NIL. The root node is the only node in the tree whose parent is NIL.
- 
- ```
graph TD; Troot[T.root] --> N1; Troot --> N2; N1 --> N3; N1 --> N4; N2 --> N5; N2 --> N6; N3 --> N7; N4 --> N8; N5 --> N9; N6 --> N10; N7 --> N11; N8 --> N12; N9 --> N13; N10 --> N14; N11 --> N15; N12 --> N16; N13 --> N17; N14 --> N18; N15 --> N19; N16 --> N20; N17 --> N21; N18 --> N22; N19 --> N23; N20 --> N24; N21 --> N25; N22 --> N26; N23 --> N27; N24 --> N28; N25 --> N29; N26 --> N30; N27 --> N31; N28 --> N32; N29 --> N33; N30 --> N34; N31 --> N35; N32 --> N36; N33 --> N37; N34 --> N38; N35 --> N39; N36 --> N40; N37 --> N41; N38 --> N42; N39 --> N43; N40 --> N44; N41 --> N45; N42 --> N46; N43 --> N47; N44 --> N48; N45 --> N49; N46 --> N50; N47 --> N51; N48 --> N52; N49 --> N53; N50 --> N54; N51 --> N55; N52 --> N56; N53 --> N57; N54 --> N58; N55 --> N59; N56 --> N60; N57 --> N61; N58 --> N62; N59 --> N63; N60 --> N64; N61 --> N65; N62 --> N66; N63 --> N67; N64 --> N68; N65 --> N69; N66 --> N70; N67 --> N71; N68 --> N72; N69 --> N73; N70 --> N74; N71 --> N75; N72 --> N76; N73 --> N77; N74 --> N78; N75 --> N79; N76 --> N80; N77 --> N81; N78 --> N82; N79 --> N83; N80 --> N84; N81 --> N85; N82 --> N86; N83 --> N87; N84 --> N88; N85 --> N89; N86 --> N90; N87 --> N91; N88 --> N92; N89 --> N93; N90 --> N94; N91 --> N95; N92 --> N96; N93 --> N97; N94 --> N98; N95 --> N99; N96 --> N100; N97 --> N101; N98 --> N102; N99 --> N103; N100 --> N104; N101 --> N105; N102 --> N106; N103 --> N107; N104 --> N108; N105 --> N109; N106 --> N110; N107 --> N111; N108 --> N112; N109 --> N113; N110 --> N114; N111 --> N115; N112 --> N116; N113 --> N117; N114 --> N118; N115 --> N119; N116 --> N120; N117 --> N121; N118 --> N122; N119 --> N123; N120 --> N124; N121 --> N125; N122 --> N126; N123 --> N127; N124 --> N128; N125 --> N129; N126 --> N130; N127 --> N131; N128 --> N132; N129 --> N133; N130 --> N134; N131 --> N135; N132 --> N136; N133 --> N137; N134 --> N138; N135 --> N139; N136 --> N140; N137 --> N141; N138 --> N142; N139 --> N143; N140 --> N144; N141 --> N145; N142 --> N146; N143 --> N147; N144 --> N148; N145 --> N149; N146 --> N150; N147 --> N151; N148 --> N152; N149 --> N153; N150 --> N154; N151 --> N155; N152 --> N156; N153 --> N157; N154 --> N158; N155 --> N159; N156 --> N160; N157 --> N161; N158 --> N162; N159 --> N163; N160 --> N164; N161 --> N165; N162 --> N166; N163 --> N167; N164 --> N168; N165 --> N169; N166 --> N170; N167 --> N171; N168 --> N172; N169 --> N173; N170 --> N174; N171 --> N175; N172 --> N176; N173 --> N177; N174 --> N178; N175 --> N179; N176 --> N180; N177 --> N181; N178 --> N182; N179 --> N183; N180 --> N184; N181 --> N185; N182 --> N186; N183 --> N187; N184 --> N188; N185 --> N189; N186 --> N190; N187 --> N191; N188 --> N192; N189 --> N193; N190 --> N194; N191 --> N195; N192 --> N196; N193 --> N197; N194 --> N198; N195 --> N199; N196 --> N200; N197 --> N201; N198 --> N202; N199 --> N203; N200 --> N204; N201 --> N205; N202 --> N206; N203 --> N207; N204 --> N208; N205 --> N209; N206 --> N210; N207 --> N211; N208 --> N212; N209 --> N213; N210 --> N214; N211 --> N215; N212 --> N216; N213 --> N217; N214 --> N218; N215 --> N219; N216 --> N220; N217 --> N221; N218 --> N222; N219 --> N223; N220 --> N224; N221 --> N225; N222 --> N226; N223 --> N227; N224 --> N228; N225 --> N229; N226 --> N230; N227 --> N231; N228 --> N232; N229 --> N233; N230 --> N234; N231 --> N235; N232 --> N236; N233 --> N237; N234 --> N238; N235 --> N239; N236 --> N240; N237 --> N241; N238 --> N242; N239 --> N243; N240 --> N244; N241 --> N245; N242 --> N246; N243 --> N247; N244 --> N248; N245 --> N249; N246 --> N250; N247 --> N251; N248 --> N252; N249 --> N253; N250 --> N254; N251 --> N255; N252 --> N256; N253 --> N257; N254 --> N258; N255 --> N259; N256 --> N260; N257 --> N261; N258 --> N262; N259 --> N263; N260 --> N264; N261 --> N265; N262 --> N266; N263 --> N267; N264 --> N268; N265 --> N269; N266 --> N270; N267 --> N271; N268 --> N272; N269 --> N273; N270 --> N274; N271 --> N275; N272 --> N276; N273 --> N277; N274 --> N278; N275 --> N279; N276 --> N280; N277 --> N281; N278 --> N282; N279 --> N283; N280 --> N284; N281 --> N285; N282 --> N286; N283 --> N287; N284 --> N288; N285 --> N289; N286 --> N290; N287 --> N291; N288 --> N292; N289 --> N293; N290 --> N294; N291 --> N295; N292 --> N296; N293 --> N297; N294 --> N298; N295 --> N299; N296 --> N300; N297 --> N301; N298 --> N302; N299 --> N303; N300 --> N304; N301 --> N305; N302 --> N306; N303 --> N307; N304 --> N308; N305 --> N309; N306 --> N310; N307 --> N311; N308 --> N312; N309 --> N313; N310 --> N314; N311 --> N315; N312 --> N316; N31
```





# BST Operations


- The search tree data structure supports many dynamic-set operations, including SEARCH, MINIMUM, MAXIMUM, PREDECESSOR, SUCCESSOR, INSERT, and DELETE.
- Thus, we can use a search tree both as a **dictionary** and as a **priority queue**.
- Basic operations on a binary search tree take time proportional to the height of the tree.
- For a complete binary tree with  $n$  nodes, such operations run in  $\Theta(\log n)$  worst-case time.



# In-order BST Walk

- The binary-search-tree property allows us to print out all the keys in a binary search tree in **sorted order** by a simple recursive algorithm, called an *inorder tree walk*.

```
INORDER-TREE-WALK( $x$ )  
1  if  $x \neq \text{NIL}$   
2      INORDER-TREE-WALK( $x.\text{left}$ )  
3      print  $x.\text{key}$   
4      INORDER-TREE-WALK( $x.\text{right}$ )
```





# BST Search (Recursive Version)

- The following procedure searches for a node with a given key in a binary search tree.
- Given a pointer to the root of the tree and a key  $k$ , TREE-SEARCH returns a pointer to a node with key  $k$  if one exists; otherwise, it returns NIL.
- The running time of TREE-SEARCH is  $O(h)$ , where  $h$  is the height of the tree.

TREE-SEARCH( $x, k$ )


```
1  if  $x == \text{NIL}$  or  $k == x.\text{key}$ 
2      return  $x$ 
3  if  $k < x.\text{key}$ 
4      return TREE-SEARCH( $x.\text{left}, k$ )
5  else return TREE-SEARCH( $x.\text{right}, k$ )
```





# BST Search Example

- To search for the key 13 in the shown tree, we follow the path  $15 \rightarrow 6 \rightarrow 7 \rightarrow 13$  from the root.





# BST Search (Iterative Version)

- We can rewrite this procedure in an iterative fashion by “unrolling” the recursion into a **while** loop. On most computers, the iterative version is more efficient

```
ITERATIVE-TREE-SEARCH( $x, k$ )  
1  while  $x \neq \text{NIL}$  and  $k \neq x.\text{key}$   
2      if  $k < x.\text{key}$   
3           $x = x.\text{left}$   
4      else  $x = x.\text{right}$   
5  return  $x$ 
```



# BST Minimum and Maximum

- The following procedures return a pointer to the minimum and maximum elements in the subtree rooted at a given node  $x$ , which we assume to be non-NIL.
- Both procedures run in  $O(h)$  time on a tree of height  $h$  since, as in TREE-SEARCH,

TREE-MINIMUM( $x$ )

```
1  while  $x.left \neq \text{NIL}$ 
2       $x = x.left$ 
3  return  $x$ 
```

TREE-MAXIMUM( $x$ )

```
1  while  $x.right \neq \text{NIL}$ 
2       $x = x.right$ 
3  return  $x$ 
```




# BST Insertion and Deletion

- The operations of insertion and deletion cause the dynamic set represented by a binary search tree to change.
- The data structure must be modified to reflect this change, but in such a way that the binary-search-tree property continues to hold.



# TREE-INSERT Example

- TREE-INSERT begins at the root of the tree and the pointer  $x$  traces a simple path downward looking for a NIL to replace with the input item  $z$ .
- The procedure maintains the **trailing pointer**  $y$  as the parent of  $x$ .
- *Example:* Inserting an item with key 13 into a binary search tree.





# BST Insertion Procedure

- The procedure inserts in binary search tree  $T$  a node  $z$  for which  $z.key = v$ ,  $z.left = \text{NIL}$ , and  $z.right = \text{NIL}$ .
- It modifies  $T$  and some of the attributes of  $z$  in such a way that it inserts  $z$  into an appropriate position in the tree.
- The procedure TREE-INSERT runs in  $O(h)$  time on a tree of height  $h$ .

TREE-INSERT( $T, z$ )

```

1   $y = \text{NIL}$ 
2   $x = T.root$ 
3  while  $x \neq \text{NIL}$ 
4       $y = x$ 
5      if  $z.key < x.key$ 
6           $x = x.left$ 
7      else  $x = x.right$ 
8   $z.p = y$ 
9  if  $y == \text{NIL}$ 
10      $T.root = z$  // tree  $T$  was empty
11 elseif  $z.key < y.key$ 
12      $y.left = z$ 
13 else  $y.right = z$ 
  
```






# BST Deletion Procedure

- The overall strategy for deleting a node  $z$  from a binary search tree  $T$  has three cases:
  1. If  $z$  has no children, then we simply remove it by modifying its parent to replace  $z$  with NIL as its child.
  2. If  $z$  has just one child, then we elevate that child to take  $z$ 's position in the tree by modifying  $z$ 's parent to replace  $z$  by  $z$ 's child.
  3. If  $z$  has two children, then we find  $z$ 's successor  $y$ , which is the minimum element in  $z$ 's right subtree, and have  $y$  take  $z$ 's position in the tree. The rest of  $z$ 's original right subtree becomes  $y$ 's new right subtree, and  $z$ 's left subtree becomes  $y$ 's new left subtree.




# BST Deletion Examples







# BST Deletion Examples (Cont'd)





# Balanced Search Trees (1 of 2)

- A binary search tree (BST) of height  $h$  can support any of the basic dynamic-set operations -such as SEARCH, PREDECESSOR, SUCCESSOR, MINIMUM, MAXIMUM, INSERT, and DELETE - in  $O(h)$  time.
  - The set operations are fast if the height of the search tree is small, which is the case if the tree is **balanced**.
  - If its height is large, however, the set operations may run no faster than with a linked list.
- To balance a BST, you need to **rebuild** it by using the node with the **median** value as its root and recursively do that with its subtrees.






# Balanced Search Trees (2 of 2)

- ***Balanced search tree*** is a search-tree data structure for which a height of  $O(\log n)$  is guaranteed.
- *Examples:* B-trees, Red-black trees, AVL trees, 2-3 trees, 2-3-4 trees.



# B-Trees

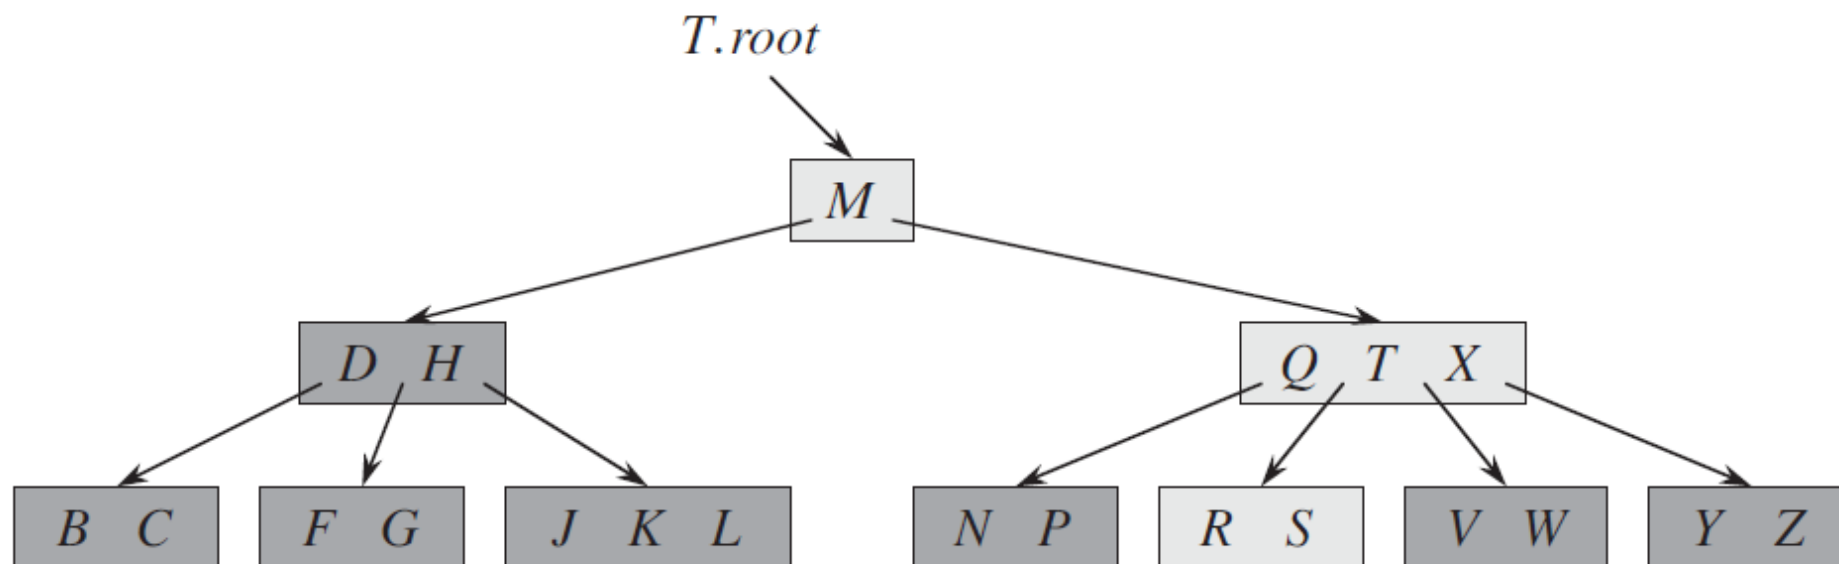
- B-trees are balanced search trees designed to work well on disks or other direct access secondary storage devices.
- B-trees nodes may have **many children**, from a few to thousands.
- Every  $n$ -node B-tree has height  **$O(\log n)$** .
- If an internal B-tree node  $x$  contains  **$x.n$**  keys, then  $x$  has  **$x.n + 1$**  children.





# B-Tree Example


- The shown B-tree has keys representing the consonants of English.
- The lightly shaded nodes are examined in a search for the letter *R*.





# Secondary Storage (1 of 2)

- Most computer systems also have **secondary storage** based on magnetic disks.
- The figure shows a typical disk drive that consists of one or more **platters**, which rotate at a constant speed around a common **spindle**.
- The drive reads and writes each platter by a **head** at the end of an **arm**.
- When a given head is stationary, the surface that passes underneath it is called a **track**.





# Secondary Storage (2 of 2)

- Although disks are cheaper and have higher capacity than main memory (RAM), they are much, much slower because they have moving mechanical parts (platter rotation and arm movement).
  - A typical RAM is **100,000 times faster** than a typical disk.
- In order to amortize the time spent waiting for mechanical movements, disks access not just one byte but several at a time.
- Information is divided into a number of equal-sized **pages** of bytes that appear consecutively within tracks, and each disk read or write is of one or more entire pages.
  - For a typical disk, a page might be  **$2^{11}$  to  $2^{14}$**  bytes in length.



# B-Tree and Secondary Storage

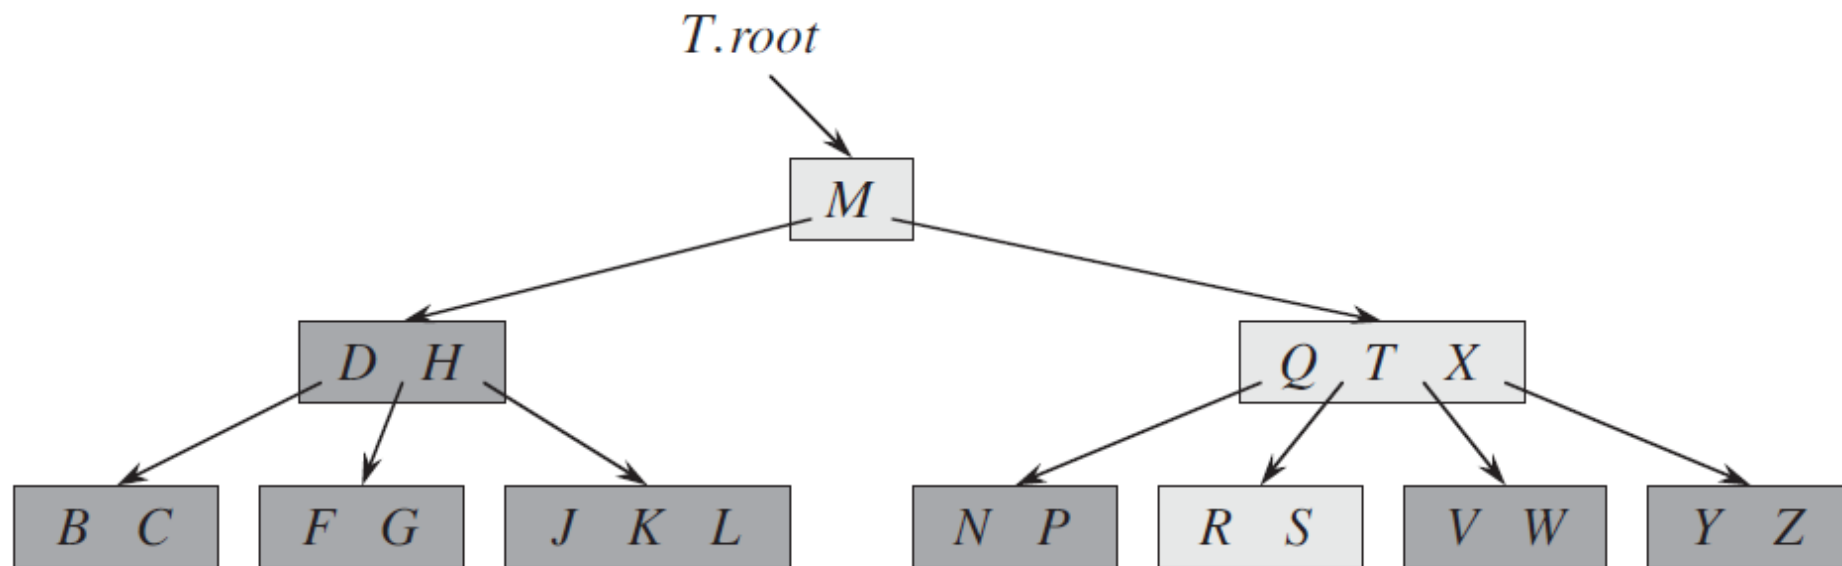
- A **B-tree node** is usually as large as a **whole page** in the disk, and this size limits the number of children a B-tree node can have.
- In a typical B-tree application, the amount of data handled is so large that all the data do not fit into main memory at once.
- The B-tree algorithms copy selected pages from disk into main memory as needed and write back onto disk the pages that have changed.
- The operation **DISK-READ( $x$ )** is used to read **node  $x$**  from the disk into main memory before we can refer to its attributes (e.g.,  $x.key_i$ ).
- The operation **DISK-WRITE( $x$ )** is used to save any changes that have been made to the attributes of node  $x$ .





# B-tree Properties (1 of 2)

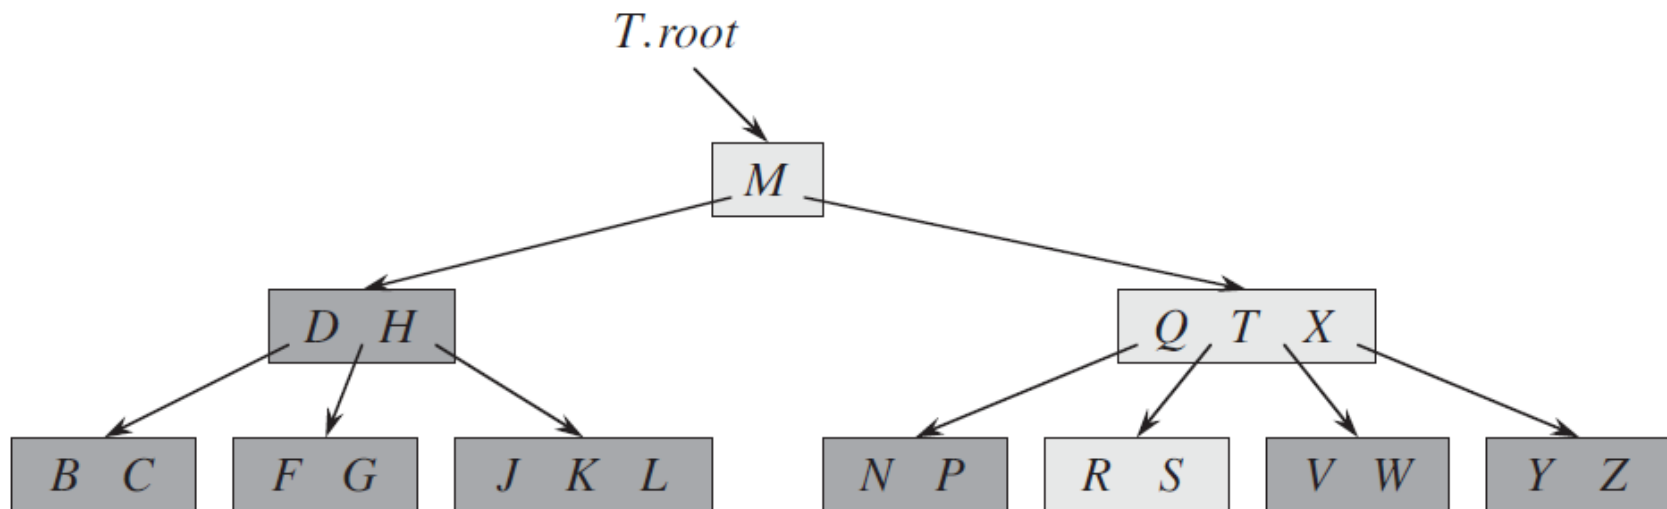
- A **B-tree**  $T$  is a rooted tree (whose root is  $T.root$ )
- Every node  $x$  has the following attributes:
  - a.  $x.n$ , the number of keys currently stored in node  $x$ ,
  - b. the  $x.n$  keys themselves are stored in increasing order, so that  $x.key_1 \leq x.key_2 \leq \dots \leq x.key_{x.n}$
  - c.  $x.leaf$ , a Boolean value that is TRUE if  $x$  is a leaf and FALSE if  $x$  is an internal node.





# B-tree Properties (2 of 2)

- Each internal (not leaf) node  $x$  also contains  $x.n + 1$  pointers  $x.c_1, x.c_2, \dots, x.c_{x.n+1}$  to its children.
  - Leaf nodes have no children, and so their  $c_i$  attributes are undefined.
- All leaves have the same depth, which is the tree's height  $h$ .
- The keys  $x.key_i$  separate the ranges of keys stored in each subtree: if  $k_i$  is any key stored in the subtree with root  $x.c_i$ , then  $k_1 \leq x.key_1 \leq k_2 \leq x.key_2 \leq \dots \leq x.key_{x.n} \leq k_{x.n+1}$





# The Minimum Degree of a B-tree

- Nodes have lower and upper bounds on the number of keys they can contain. We express these bounds in terms of a fixed integer  $t \geq 2$  called the ***minimum degree*** of the B-tree:
  - a. Every node other than the root must have at least  **$t - 1$  keys**.
  - b. An internal node (i.e., neither a root nor a leaf) with the least number of keys thus has  **$t$  children**.
  - c. If the tree is nonempty, the root must have at least **one key**.
  - d. Every node may contain at most  **$2t - 1$  keys**.
  - e. A non-leaf node with the most number of keys thus has  **$2t$  children**.
  - f. We say that a node is ***full*** if it contains exactly  **$2t - 1$  keys**.



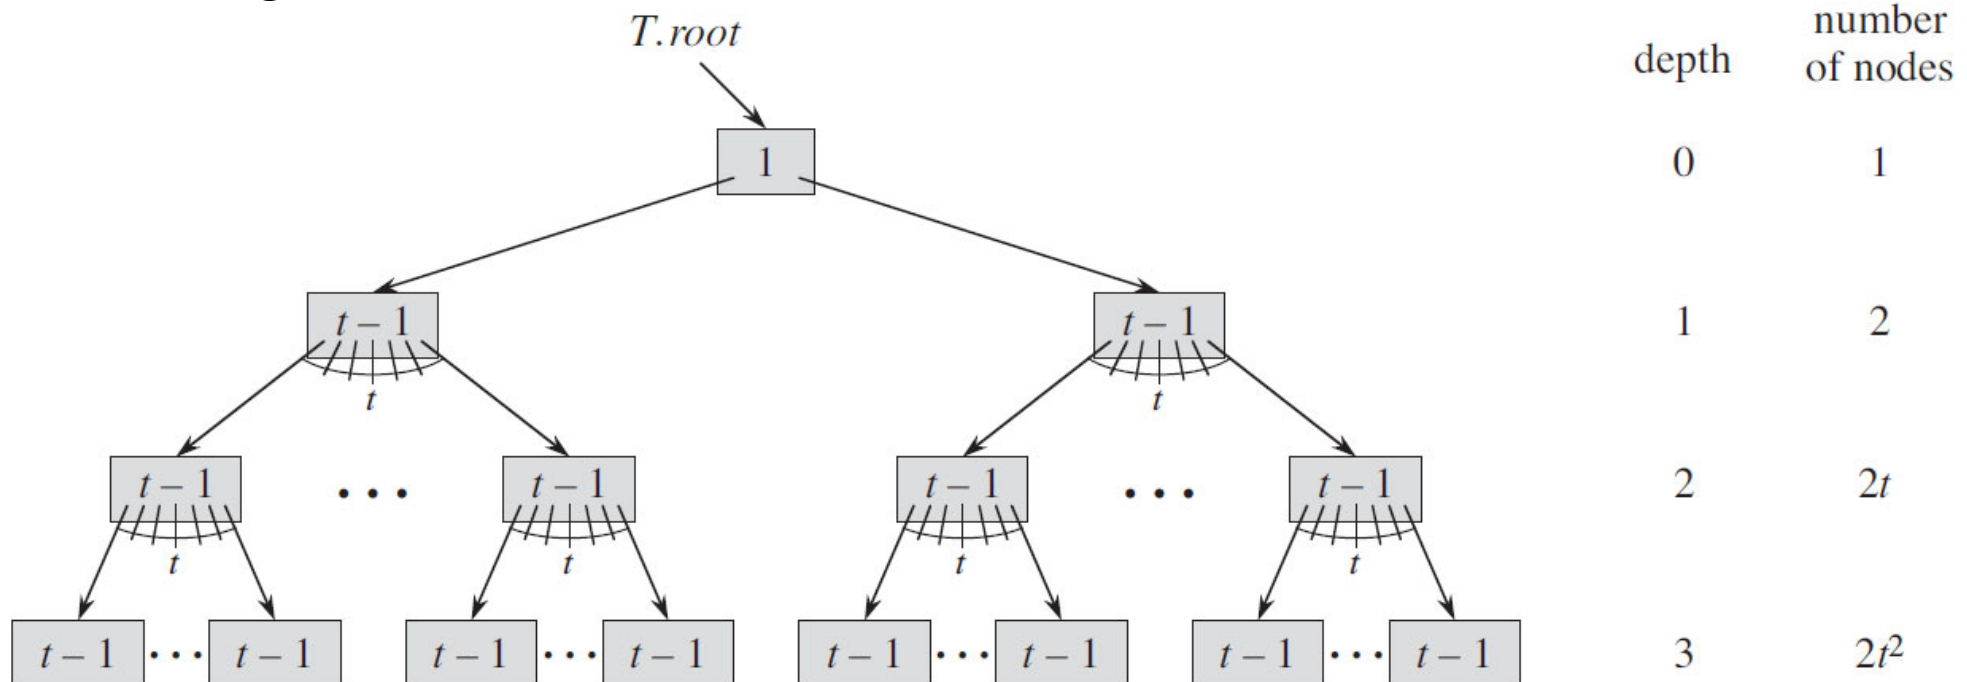
# The Simplest B-tree

- The simplest B-tree occurs when  $t = 2$ , and internal nodes have from  $t-1$  to  $2t-1$  keys.
- Every internal node then has either 2, 3, or 4 children, and we call this tree a **2-3-4 tree**.
- In practice, however, much larger values of  $t$  yield B-trees with smaller height.



# The Height of a B-tree

- The number of disk accesses for most operations on a B-tree is proportional to the height of the B-tree. What is the **worst-case height** of a B-tree? Worst case happens if every node has the minimum allowed number of keys.
- Theorem:** If  $n \geq 1$ , then for any  $n$ -node B-tree  $T$  of height  $h$  and minimum degree  $t \geq 2 \rightarrow h \leq \log_t \frac{n+1}{2}$
- The figure illustrates such a tree for  $h = 3$ .



$$n-1 = 2+2t+2t^2+\dots +2t^{h-1} = 2 \frac{(t^h-1)}{(t-1)} \quad (\text{for max } h, t=2 \rightarrow (n+1)/2 = t^h)$$



# Basic Operations on B-trees

---

- Basic operations on B-trees include:
  - B-TREE-CREATE
  - B-TREE-SEARCH
  - B-TREE-INSERT
  - B-TREE-DELETE



# Creating an Empty B-Tree

- To build a B-tree  $T$ , we first use B-TREE-CREATE to create an empty root node and then call B-TREE-INSERT to add new keys.
- Both of these procedures use an auxiliary procedure ALLOCATE-NODE, which allocates one disk page to be used as a new node in  $O(1)$  time.
  - Note here the address of a node is not a pointer to a main memory location, rather an address in the disk.

B-TREE-CREATE( $T$ )

```
1   $x = \text{ALLOCATE-NODE}()$ 
2   $x.\text{leaf} = \text{TRUE}$ 
3   $x.n = 0$ 
4   $\text{DISK-WRITE}(x)$ 
5   $T.\text{root} = x$ 
```



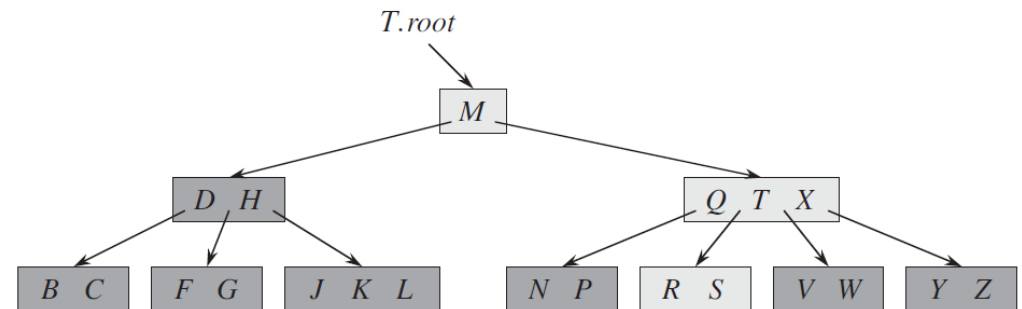
# Searching a B-tree

- To search a B-tree instead of making a two-way branching decision at each node (as in binary search), we make at each internal node  $x$ , an  $x.n+1$  way **branching** decision.
- Lines 1–3 find the smallest index  $i$  such that  $k \leq x.key_i$ , or else they set  $i$  to  $x.n + 1$ .
- The return value here is the address of the B-Tree node where  $k$  is found and its index within the node.
- The procedure accesses  $O(h) = O(\log_t n)$  disk pages, where  $h$  is the height of the B-tree and  $n$  is the number of nodes in the B-tree

**B-TREE-SEARCH**( $x, k$ )

```

1   $i = 1$ 
2  while  $i \leq x.n$  and  $k > x.key_i$ 
3       $i = i + 1$ 
4  if  $i \leq x.n$  and  $k == x.key_i$ 
5      return ( $x, i$ )
6  elseif  $x.leaf$ 
7      return NIL
8  else DISK-READ( $x.c_i$ )
9      return B-TREE-SEARCH( $x.c_i, k$ )
  
```








# Inserting a Key into a B-Tree

- With a B-tree, however, we cannot simply create a new leaf node and insert it, as the resulting tree would fail to be a valid B-tree.
- We insert the new key into an existing **leaf** node.
- Since we cannot insert a key into a leaf node that is full, we introduce an operation that **splits** a full node  $y$  (having  $2t - 1$  keys) around its **median key**  $y.key_t$  into two nodes having only  $t - 1$  keys each.
- The median key moves up into  $y$ 's parent to identify the dividing point between the two new trees. But if  $y$ 's parent is also full, we must split it before we can insert the new key.
- To split a full root, we will first make the root a child of a new empty root node, so that we can use B-TREE-SPLIT-CHILD.
  - The tree thus grows in height by one; splitting is the only means by which the tree grows.

# Example of Splitting a Node


- Splitting a node with  $t = 4$ . Node  $y = x.c_i$  splits into two nodes,  $y$  and  $z$ , and the median key  $S$  of  $y$  moves up into  $y$ 's parent





# Example of Splitting the Root




- Splitting the root with  $t = 4$ .
- Root node  $r$  splits in two.
- A new root node  $s$  is created. The new root contains the median key of  $r$  and has the two halves of  $r$  as children.
- The B-tree grows in height by one when the root is split.
  - Unlike a binary search tree, a B-tree increases in height at the top instead of at the bottom.





# Example of Inserting a Key (1 of 2)


- The minimum degree  $t$  for this B-tree is 3, so a node can hold at most 5 keys and at least 2 keys. A key is inserted in an existing leaf node



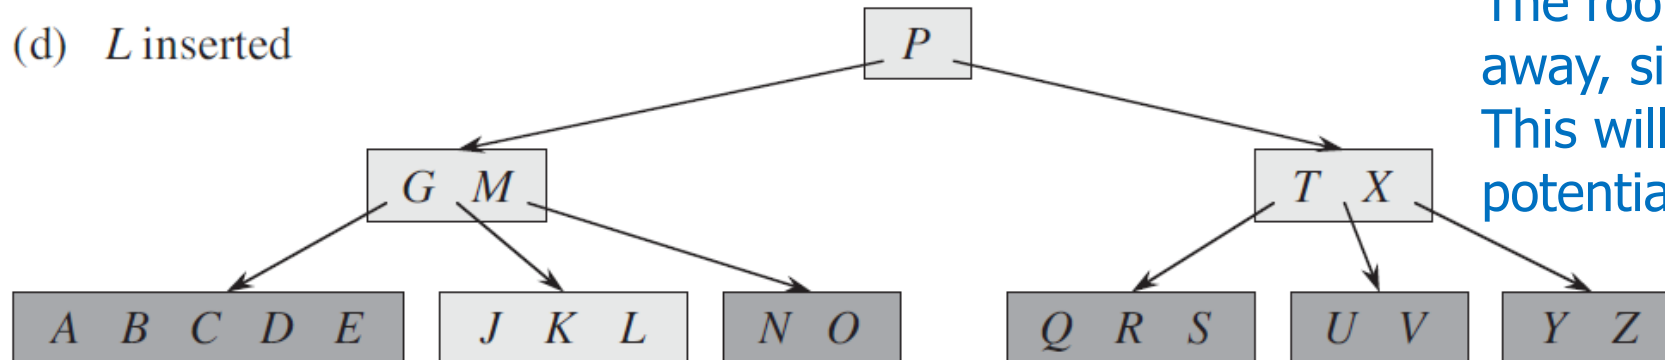


# Example of Inserting a Key (2 of 2)

(c) *Q* inserted




(d) *L* inserted



The root splits right away, since it is full. This will avoid any potential backtracking.

(e) *F* inserted






# Deleting a Key from a B-Tree



- As in insertion, we must guard against deletion producing a tree whose structure violates the B-tree properties.
  - Just as we had to ensure that a node didn't get too big due to insertion, we must ensure that a node doesn't get too small during deletion (except that the root is allowed to have fewer than the minimum number  $t - 1$  of keys).
- Deletion from a B-tree is a little more complicated than insertion, because we can delete a key from any node—not just a leaf—and when we delete a key from an internal node, we will have to rearrange the node's children.

# Example of Deleting a Key (1 of 2)

(a) initial tree



➤ The minimum degree  $t$  for this B-tree is 3, so a node (other than the root) cannot have fewer than 2 keys.


(b)  $F$  deleted:(c)  $M$  deleted:






# Example of Deleting a Key (2 of 2)


(c) *M* deleted:




(d) *G* deleted:



(e) *D* deleted:



(f) *B* deleted:



Even merge here is not required, but it is a good practice to merge nodes with  $t-1$  keys to avoid a potential backtracking as sometimes a key may have to be moved into a child node as with deleting key *B* here.