


EECE7205: Fundamentals of Computer Engineering




Graph Algorithms



Directed Graph


- A **directed graph** (or **digraph**) G is a pair (V, E) , where V is a finite set and E is a binary relation on V .
- The set V is called the **vertex set** of G , and its elements are called **vertices** (singular: **vertex**).
- The set E is called the **edge set** of G , and its elements are called **edges**.
- Vertices are represented by circles, and edges are represented by arrows.
- In the shown example:
 $V = \{1, 2, 3, 4, 5, 6\}$
 $E = \{(1, 2), (2, 2), (2, 4), (2, 5), (4, 1), (4, 5), (5, 4), (6, 3)\}.$
- The edge $(2, 2)$ is a self-loop.





Undirected Graph

- In an ***undirected graph*** $G = (V, E)$, the edge set E consists of *unordered* pairs of vertices, rather than ordered pairs.
- An edge is a set $\{u, v\}$ where $u, v \in V$ and $u \neq v$.
- In the shown example:
 $V = \{1, 2, 3, 4, 5, 6\}$
 $E = \{(1, 2), (1, 5), (2, 5), (3, 6)\}$.
- In an undirected graph, **self-loops** and **multiple edges between the same pair of nodes** are not allowed.





Graph Definitions (1 of 3)

- In a directed graph, an edge (u, v) is ***incident from*** or ***leaves*** vertex u and is ***incident to*** or ***enters*** vertex v .
- In an undirected graph, an edge (u, v) is ***incident on*** vertices u and v .
- In a graph with an edge (u, v) means that vertex v is ***adjacent*** to vertex u . In an undirected graph, the adjacency relation is symmetric.
- The ***degree*** of a vertex in an undirected graph is the number of edges incident on it. A vertex with degree 0 is called ***isolated***.
- In a directed graph, the ***out-degree*** of a vertex is the number of edges leaving it, and the ***in-degree*** of a vertex is the number of edges entering it.
 - The ***degree*** of a vertex in a directed graph is its in-degree plus its out-degree.



Graph Definitions (2 of 3)

- A **path** of **length** k from a vertex v_0 to a vertex v_k in a graph is the sequence of vertices $(v_0, v_1, v_2, \dots, v_k)$ where $(v_{i-1}, v_i) \in E$ for $i = 1, 2, \dots, k$.
- The length of the path is the number of edges in the path.
- If there is a path p from u to v , we say that v is **reachable** from u via p .
- A path is **simple** if all vertices in the path are distinct.
- A **subpath** of path $p = (v_0, v_1, v_2, \dots, v_k)$ is a contiguous subsequence of its vertices.
- In a directed graph, a path $(v_0, v_1, v_2, \dots, v_k)$ forms a **cycle** if $v_0 = v_k$ and it is a **simple cycle** if all its other vertices are distinct.
- A self-loop is a cycle of length 1. A directed graph with no self-loops is **simple**.
- In an undirected graph, a path $(v_0, v_1, v_2, \dots, v_k)$ forms a **cycle** if $k \geq 3$ and $v_0 = v_k$ and it is a **simple cycle** if all its vertices are distinct.
- A graph with no cycles is **acyclic**.




Graph Definitions (3 of 3)

- An undirected graph is **connected** if every vertex is reachable from all other vertices.
- A directed graph is **strongly connected** if every two vertices are reachable from each other.
- A **complete graph** is an undirected graph in which every pair of vertices is adjacent.
- An acyclic, undirected graph is a **forest**.
- A connected, acyclic, undirected graph is a **tree**.
- A **Sparse** graph is a graph with $|E|$ is much less than $|V|^2$.
- A **Dense** graph is a graph with $|E|$ is close to $|V|^2$.




Representing Undirected Graphs

- The adjacency-list representation is usually the method of choice for **sparse** graphs.
- The adjacency-matrix representation is usually the method of choice for **dense** graphs or when we need to be able to tell quickly if there is an edge connecting two given vertices.
- The sum of the lengths of all the adjacency lists is $2|E|$
- Observe the symmetry along the main diagonal of the adjacency matrix.



An undirected graph G



An adjacency-list representation of G


	1	2	3	4	5
1	0	1	0	0	1
2	1	0	1	1	1
3	0	1	0	1	0
4	0	1	1	0	1
5	1	1	0	1	0

An adjacency-matrix representation of G




Representing Directed Graphs

- The sum of the lengths of all the adjacency lists is $|E|$



A directed graph G



An adjacency-list representation of G

	1	2	3	4	5	6
1	0	1	0	1	0	0
2	0	0	0	0	1	0
3	0	0	0	0	1	1
4	0	1	0	0	0	0
5	0	0	0	1	0	0
6	0	0	0	0	0	1

An adjacency-matrix representation of G




Graph Traversals


- Given a graph $G = (V, E)$ and a distinguished **source** vertex s , graph traversal means systematically explores the edges of G to “discover” every vertex that is reachable from s .
 - Graph traversal also answers the question: **is there a path from s to t ?**
- There are two approaches of graph traversal:
 - **Breadth-First Search (BFS)** discovers all vertices at distance k from s before discovering any vertices at distance $k + 1$.
 - **Depth-First Search (DFS)** searches “deeper” in the graph whenever possible before backing up.
- Unlike a tree traversal, which always visits all of the nodes in a tree, a graph traversal does not necessarily visit all of the vertices in the graph unless the graph is connected.
 - If a graph traversal does not visit all vertices in the graph, then the graph is **not connected**.



BFS vs DFS



BFS Example



DFS Example



Breadth-First Search (1 of 2)

- BFS strategy: discovers all vertices at distance k from s before discovering any vertices at distance $k + 1$
- This strategy is reflected in using a **queue** as in the following algorithm:

```


bfs(v: Vertex)
  q = new empty queue
  q.enqueue(v) // Add v to queue
  Mark v as visited
  while (!q.isEmpty()) {
    q.dequeue(w)
    for (each unvisited vertex u adjacent to w) {
      Mark u as visited
      q.enqueue(u)
    }
  }

```



Breadth-First Search (2 of 2)

- The results of a breadth-first traversal, beginning at vertex *a*, of the following graph, is: *a, b, f, i, c, e, g, d, h*



```

bfs(v: Vertex)
  q = a new empty queue
  q.enqueue(v) // Add v to queue
  Mark v as visited
  while (!q.isEmpty()) {
    q.dequeue(w)
    for (each unvisited vertex u adjacent to w) {
      Mark u as visited
      q.enqueue(u) }
  }

```

Node visited	Queue (front to back)
a	a (empty)
b	b
f	b f
i	b f i
c	f i
e	f i c
g	f i c e
d	i c e
h	i c e g
	c e g
	e g
	e g d
	g d
	d
	(empty)
	h
	(empty)



Depth-First Search (1 of 2)


- DFS strategy: searches “deeper” in the graph whenever possible before backing up.
- This strategy is reflected in using a **stack** as in the following algorithm:

```
dfs(v: Vertex)
  s = a new empty stack
  s.push(v) // Push v onto the stack
  Mark v as visited
  while (!s.isEmpty()) {
    if (no unvisited vertices are adjacent to
        the vertex on the top of the stack)
      s.pop() // Backtrack
    else {
      Select an unvisited vertex u adjacent to the
        vertex on the top of the stack
      s.push(u)
      Mark u as visited
    }
  }
```



Depth-First Search (2 of 2)

- The results of a depth-first traversal, beginning at vertex *a*, of the following graph, is: *a, b, c, d, g, e, f, h, i*



```
dfs(v: Vertex)
  s = a new empty stack
  s.push(v) // Push v onto the stack
  Mark v as visited
  while (!s.isEmpty()) {
    if (no unvisited vertices adjacent to top vertex)
      s.pop() // Backtrack
    else {
      Select an unvisited vertex u adjacent to top vertex
      s.push(u)
      Mark u as visited }
  }
```

<u>Node visited</u>	<u>Stack (bottom to top)</u>
a	a
b	a b
c	a b c
d	a b c d
g	a b c d g
e	a b c d g e
(backtrack)	a b c d g
f	a b c d g f
(backtrack)	a b c d g
(backtrack)	a b c d
h	a b c d h
(backtrack)	a b c d
(backtrack)	a b c
(backtrack)	a b
(backtrack)	a
i	a i
(backtrack)	a
(backtrack)	(empty)



Detect Cycles in a Graph (1 of 3)

- Recall that a connected undirected graph cannot contain a cycle if it has n vertices and exactly $n - 1$ edges,.
- The following is a modified DFS algorithm to determine whether a general graph contains a cycle.
- In this algorithm, each vertex has the following attributes:
 - **visited**: to indicate whether the vertex has been visited.
 - **predecessor**: to mark its predecessor vertex on the DFS path.
 - **onStack**: to indicate whether the vertex is still in the stack.
 - **finished**: to indicate that the vertex is not in the stack anymore.
- If the vertex at the top of the stack has a neighbor that is still in the stack and it is not its predecessor, then a cycle exists.



Detect Cycles in a Graph (2 of 3)


```

hasCycle(v: Vertex): Boolean {
    s = a new empty stack
    Mark all vertices as unvisited and not onStack
    s.push(v)
    v.onStack = true
    while (!s.isEmpty()) {
        top = s.peek()
        if (top has a neighbor marked as onStack and
            it is not top's predecessor)
            return true // A cycle exists
        else if (top has an unvisited neighbor u) {
            u.onStack=true
            u.predecessor = top
            s.push(u)
        }
        else { top.finished=true
                s.pop()}
    }
    return false
}
  
```




Detect Cycles in a Graph (3 of 3)


- Applying the `hasCycle` algorithm on the following graph, starting from vertex *a*, will result in a stack with the following contents (from bottom to top): *a*, *b*, *c*, *d*, *g*, *e*
- As the top of the stack, vertex *e*, has a neighbor marked as `onStack` and is not *e*'s predecessor (vertex *b*), then the algorithm will return `true`.






BFS and DFS Trees

- The predecessor of each vertex in the BFS and DFS can be maintained as explained in the detect cycle algorithm.
- The BFS and DFS predecessor subgraph forms a tree as shown.



BFS Predecessor Subgraph
Tree rooted at vertex *a*



DFS Predecessor Subgraph
Tree rooted at vertex *a*




Printing BFS and DFS Paths

- The following procedure prints out the vertices on a path from vertex s to v , on the BFS and DFS predecessor subgraph.
- Assume all predecessors are initialized to NIL

```

PrintPath(G: Graph, s: Vertex, v: Vertex)
  if v == s
    print s
  else if v.predecessor == NIL
    print no path from s to v exists
  else PrintPath(G, s, v.predecessor)
    print v
  }
  
```





Spanning Trees


- Application example:

Electronic circuit designs often need to make the pins of several components electrically equivalent by wiring them together. To interconnect a set of n pins, we can use an arrangement of $n - 1$ wires, each connecting two pins. Of all such arrangements, the one that uses the least amount of wire is usually the most desirable.



Minimum Spanning Trees


- We can model the wiring problem with a connected, undirected graph $G = (V, E)$, where V is the set of pins, E is the set of possible interconnections between pairs of pins.
- For each edge $(u, v) \in E$, we have a weight $w(u, v)$ specifying the cost (amount of wire needed) to connect u and v .
- We then wish to find an acyclic subset $T \subseteq E$ that connects all of the vertices and whose total weight is minimized.





Spanning Tree Example

- Since T is acyclic and connects all of the vertices, it must form a tree, which we call a ***spanning tree*** since it “spans” the graph G .
- We call the problem of determining the tree T the ***minimum-spanning-tree problem*** (MST).





MST: Kruskal's Algorithm

MST-KRUSKAL(G, w)


```
1   $A = \emptyset$ 
2  for each vertex  $v \in G.V$ 
3      MAKE-SET( $v$ )
4  sort the edges of  $G.E$  into nondecreasing order by weight  $w$ 
5  for each edge  $(u, v) \in G.E$ , taken in nondecreasing order by weight
6      if FIND-SET( $u$ )  $\neq$  FIND-SET( $v$ )
7           $A = A \cup \{(u, v)\}$ 
8          UNION( $u, v$ )
9  return  $A$ 
```









To combine two sets
to form one tree

Comparing u 's set with v 's
set to determine whether
these vertices u and v belong
to the same tree






Kruskal's Algorithm Example











MST: Prim's Algorithm

- Prim's algorithm has the property that the edges in the set A always form a single tree.
- The tree starts from an arbitrary root vertex r and grows until the tree spans all the vertices in V .
- Each step adds to the tree A the "lightest" edge that connects A to an isolated vertex—one on which no edge of A is incident.
- Example:






Prim's Algorithm Example (Cont'd)





MST Exercise

- Find the MST of the following graph using both Kruskal's Algorithm and Prim's Algorithm.



graph edges
sorted by weight

↓

0-7	16
2-3	17
1-7	19
0-2	26
5-7	28
1-3	29
1-5	32
2-7	34
4-5	35
1-2	36
4-7	37
0-4	38
6-2	40
3-6	52
6-0	58
6-4	93




Shortest Paths

- Application example:

Finding the shortest possible route from *Boston* to *Providence*.

Given a road map of the *United States* on which the distance between each pair of adjacent intersections is marked.





Shortest Paths


- Consider a weighted, directed graph $G = (V, E)$ with edge-weight function : $w : E \rightarrow \mathbb{R}$.
- The **weight** of path $p = v_0 \rightarrow v_1 \rightarrow \cdots \rightarrow v_k$ is defined to be

$$w(p) = \sum_{i=1}^k w(v_{i-1}, v_i) .$$

- A **shortest path** from u to v is a path of minimum weight from u to v . The **shortest-path weight** from u to v is defined as

$$\delta(u, v) = \min\{w(p) : p \text{ is a path from } u \text{ to } v\}.$$


- **Note:** $\delta(u, v) = \infty$ if no path from u to v exists.





Edge Weights


- Edge weights can represent metrics such as distances, time, cost, penalties, loss, or any other quantity that accumulates linearly along a path and that we would want to minimize.
- The breadth-first-search (BFS) algorithm, as shown below, is a shortest-paths algorithm that works on unweighted graphs, that is, graphs in which each edge has unit weight.





Single-Source Shortest Paths (1 of 2)




- The **single-source shortest-paths problem**: given a graph $G = (V, E)$, we want to find a shortest path from a given **source** vertex $s \in V$ to each vertex $v \in V$.
- The algorithm for the single-source problem can solve many other problems, including the following variant:
 - **Single-destination shortest-paths problem**: Find a shortest path to a given **destination** vertex t from each vertex v . By reversing the direction of each edge in the graph and then solve the single-source problem.





Single-Source Shortest Paths (2 of 2)

- Shortest paths are not necessarily unique.
- Example:* The following figure shows a weighted, directed graph and two shortest-paths trees with the same root.






Initialization

- The shortest paths algorithms start with the initialization process where for each vertex $v \in V$:
 - $v.d$ is the upper bound on the weight of a shortest path from source s to v and it is initialized to ∞ .
 - $v.d$ is called a ***shortest-path estimate***.
 - $v.\pi$ is vertex v 's ***predecessor***.

INITIALIZE-SINGLE-SOURCE(G, s)

```

1  for each vertex  $v \in G.V$ 
2       $v.d = \infty$ 
3       $v.\pi = \text{NIL}$ 
4   $s.d = 0$ 
  
```





Relaxation (1 of 2)



- The process of **relaxing** an edge (u, v) consists of testing whether we can improve the shortest path to v that has been found so far by going through u and, if so, updating $v.d$ and $v.\pi$.

RELAX(u, v, w)

```

1  if  $v.d > u.d + w(u, v)$ 
2       $v.d = u.d + w(u, v)$ 
3       $v.\pi = u$ 
  
```

- Example:*





Relaxation (2 of 2)

- Shortest path algorithms calls INITIALIZE-SINGLE-SOURCE once and then repeatedly relaxes edges.
- Relaxation is the only means by which shortest path estimates and predecessors change.
- The algorithms differ in how many times they relax each edge and the order in which they relax edges.



Dijkstra's Algorithm (1 of 2)

- Dijkstra's algorithm solves the single-source shortest-paths problem on a weighted, directed graph $G = (V, E)$ for the case in which all edge weights are **nonnegative**.
- Dijkstra's algorithm maintains a **set S** of vertices whose final shortest-path weights from the source s have already been determined.
- The algorithm repeatedly selects the vertex $u \in V - S$ with the minimum shortest-path estimate, adds u to S , and relaxes all edges leaving u .
- The algorithm uses a **min-priority queue Q** of vertices, keyed by their d values (the shortest-path estimate).



Dijkstra's Algorithm (2 of 2)


DIJKSTRA(G, w, s)

```
1  INITIALIZE-SINGLE-SOURCE( $G, s$ )
2   $S = \emptyset$ 
3   $Q = G.V$ 
4  while  $Q \neq \emptyset$ 
5       $u = \text{EXTRACT-MIN}(Q)$ 
6       $S = S \cup \{u\}$ 
7      for each vertex  $v \in G.Adj[u]$ 
8          RELAX( $u, v, w$ )
```




Dijkstra's Algorithm Example


- Black vertices are in the set S , and white vertices are in the min-priority queue $Q = V - S$. The shaded vertex is vertex u with the minimum d value.




(a)




(b)




(c)



(d)



(e)




(f)



Dijkstra's Algorithm Exercise

- Find the single-source shortest-paths from vertex s in the shown weighted, directed graph.







Negative-Weight Edges

- Some instances of the single-source shortest-paths problem may include edges whose weights are negative.
- If the graph $G = (V, E)$ contains no negative-weight cycles reachable from the source s , then for all $v \in V$, the shortest-path weight $\delta(s, v)$ remains well defined, even if it has a negative value.
- If the graph contains a negative-weight cycle reachable from s , however, shortest-path weights are not well defined. No path from s to a vertex on the cycle can be a shortest path.
- If there is a negative-weight cycle on some path from s to v , we define $\delta(s, v) = -\infty$.



Negative-Weight Cycle Example

- In the following example, The shortest-path weight from source s appears within each vertex.
- Vertices e , f , and g are reachable from s through a negative-weight cycle (cycle e, f), they have shortest-path weights of $-\infty$.
- Vertices h , i , and j are not reachable from s , and so their shortest-path weights are ∞ , even though they lie on a negative-weight cycle.





Properties of Shortest Paths

■ Triangle inequality

- For any edge $(u, v) \in E$, we have $\delta(s, v) \leq \delta(s, u) + w(u, v)$.

■ Upper-bound property


- We always have $v.d \geq \delta(s, v)$ for all vertices $v \in V$, and once $v.d$ achieves the value $\delta(s, v)$, it never changes.

■ No-path property

- If there is no path from s to v , then we always have $v.d = \delta(s, v) = \infty$.

■ Predecessor-subgraph property

- Once $v.d = \delta(s, v)$ for all $v \in V$, the predecessor subgraph is a shortest-paths tree rooted at s .





The Bellman-Ford Algorithm (1 of 3)

- The ***Bellman-Ford algorithm*** solves the single-source shortest-paths problem in the general case in which edge weights may be negative.
- It returns a Boolean value indicating whether or not there is a negative-weight cycle that is reachable from the source.
 - If there is such a cycle, the algorithm indicates that no solution exists.
 - If there is no such cycle, the algorithm produces the shortest paths and their weights.



The Bellman-Ford Algorithm (2 of 3)

BELLMAN-FORD(G, w, s)

```

1  INITIALIZE-SINGLE-SOURCE( $G, s$ )
2  for  $i = 1$  to  $|G.V| - 1$ 
3      for each edge  $(u, v) \in G.E$ 
4          RELAX( $u, v, w$ )
5  for each edge  $(u, v) \in G.E$ 
6      if  $v.d > u.d + w(u, v)$ 
7          return FALSE
8  return TRUE
  
```

The algorithm makes $|V| - 1$ passes over the edges of the graph. Each pass consists of relaxing each edge of the graph once.

It checks for a negative-weight cycle and returns the appropriate Boolean value.

Triangle inequality




The Bellman-Ford Algorithm (3 of 3)


- The previous algorithm can be optimized by:
 - Maintaining an array D with entries *True* or *False* for whether the weight of each vertex has been changed or not (initially this array contain all *False* except a *True* entry for source node s)
 - Skipping the $Relax(u, v, w)$ step if u has entry *False* in array D .
 - After each call to $Relax(u, v, w)$, update v 's entry in D to indicate whether v 's weight has been changed or not.
 - No need to continue the first *for* loop if all entries in D are false.
- This optimization of the algorithm uses the dynamic programming approach (to be discussed later in the course).




Bellman-Ford Example (1 of 2)



(a)




(b)




(c)

Bellman-Ford Example (2 of 2)



(d)




(e)

- As the triangle inequality ($\delta(s, v) \leq \delta(s, u) + w(u, v)$) can be verified on all edges, then the Bellman-Ford algorithm returns TRUE in this example.



Dijkstra's vs. Bellman-Ford

- Use both Dijkstra's and Bellman-Ford to solve the single-source shortest-paths problem for the shown graph where node s is the source.
- Using Dijkstra's will not produce the correct answer while Bellman-Ford will result in the correct answer.
- So, Dijkstra's does not work when edges have negative weight regardless if there is a negative cycle or not.






Flow Network

- We modelled a road map as a directed graph in order to find the shortest path from one point to another, we can also interpret a directed graph as a “flow network” and use it to answer questions about material flows.
- Imagine a material coursing through a system from a source, where the material is produced, to a sink, where it is consumed.
- The source produces the material at some steady rate, and the sink consumes the material at the same rate.
- The “flow” of the material at any point in the system is intuitively the rate at which the material moves.
- Flow networks can model many problems, including liquids flowing through pipes, parts through assembly lines, current through electrical networks, and information through communication networks.




Flow Networks Properties

- A **flow network** $G = (V, E)$ is a directed graph in which each edge $(u, v) \in E$ has a nonnegative **capacity** $c(u, v)$.
- It is required that **self-loops are not allowed**.
- If E contains an edge (u, v) then there is **no edge (v, u) in the reverse direction**.
- Two vertices are distinguished in a flow network: a **source** s and a **sink** t .
- *Assumption:* each vertex lies on some path from the source to the sink. The graph is therefore **connected** and each vertex other than s has at least one entering edge.





Flow Definition



- Let $G = (V, E)$ be a flow network with a capacity function \mathbf{c} and let \mathbf{s} be the source of the network, and let \mathbf{t} be the sink. A **flow** in G is a real-valued function $f: V \times V \rightarrow \mathbb{R}$ that satisfies the following two properties:

- Capacity constraint:** For all $u, v \in V$, we require

$$0 \leq f(u, v) \leq c(u, v)$$

which says that the flow from one vertex to another must be nonnegative and must not exceed the given capacity

- Flow conservation:** For all $u \in V - \{s, t\}$, we require “*flow in equals flow out*”

$$\sum_{v \in V} f(v, u) = \sum_{v \in V} f(u, v) .$$

When $(u, v) \notin E$, there can be no flow from u to v , and $f(u, v)=0$.



The Maximum Flow Problem

- In the maximum-flow problem, we wish to compute the greatest rate at which we can ship material from the source to the sink without violating any capacity constraints.
- The total flow, f , out of the source s , is defined as:



$$|f| = \sum_{v \in V} f(s, v)$$

- In the **maximum-flow problem**, we are given a flow network G with source s and sink t , and we wish to find a flow f of maximum value.



Flow Network Example (1 of 2)



- The left figure model the trucking problem of a company *LP*. *LP* has a factory (source *s*) in Vancouver that manufactures furniture, and it has a warehouse (sink *t*) in Winnipeg that stocks them.
- The trucks ships the furniture over specified routes (edges) between cities (vertices). Routes have a limited capacity of at most $c(u, v)$ furniture containers per day between each pair of cities *u* and *v*.
- *LP* needs to determine the largest number *p* of containers per day that they can ship. *LP* is not concerned with how long it takes for the furniture to get from the factory to the warehouse; they care only in maximizing *p*.
- The right figure represent one possible solution. Is it maximum?





Flow Network Example (2 of 2)

- The top solution has a flow of 19, while the bottom solution has a better flow of 23.
- But how can we improve the top solution to reach the bottom one?
- As you see, we had to remove an already assigned flows (v_2 to v_1 and v_3 to v_2).







Residual Networks

- The figures show the flow network G followed by its residual network G_f .
- Edges in G_f have the following "residual capacity" c_f in any edge (u, v)

$$c_f(u, v) = \begin{cases} c(u, v) - f(u, v) & \text{if } (u, v) \in E, \\ f(v, u) & \text{if } (v, u) \in E, \\ 0 & \text{otherwise.} \end{cases}$$

- Edges with $c_f = 0$ are not included in G_f .
- Observe that the residual network G_f does not satisfy the definition of a flow network because it may contain both an edge (u, v) and its reversal (v, u) .





Augmentation

- We define $f \uparrow f'$, the **augmentation** of flow f by f' , to be:


$$(f \uparrow f')(u, v) = \begin{cases} f(u, v) + f'(u, v) - f'(v, u) & \text{if } (u, v) \in E, \\ 0 & \text{otherwise.} \end{cases}$$

- The flow on (u, v) is increased by $f'(u, v)$ but is decreased by $f'(v, u)$ because pushing flow on the reverse edge in the residual network signifies decreasing the flow in the original network.
- Pushing flow on the reverse edge in the residual network is also known as **cancellation**.



Augmenting Paths

- Given a residual network $G_f = (V, E_f)$ and a flow f , an *augmenting path* p is a simple path from s to t in G_f .
- The shaded path in the shown figure is an augmenting path.
- Since the smallest residual capacity on this path is $c_f(v_2, v_3) = 4$, the flow through each edge of this path can be increased by up to 4 units without violating a capacity constraint.





Residual Capacity

- The ***residual capacity*** of an augmenting path p is the maximum amount by which we can increase the flow on each edge in p and it is given by:

$$c_f(p) = \min \{c_f(u, v) : (u, v) \text{ is on } p\}$$


- Leading to the following definition:

$$f_p(u, v) = \begin{cases} c_f(p) & \text{if } (u, v) \text{ is on } p, \\ 0 & \text{otherwise.} \end{cases}$$


- **Corollary:** If we augment the current flow f by f_p , we get another flow in G whose value is closer to the maximum.




Augmentation Example




G with flow f



G_f with flow $f_p = 4$



G after augmenting f with f_p



Updated G_f



Cuts of Flow Networks

- A **cut** (S, T) of flow network $G = (V, E)$ is a partition of V into S and T where $T = V - S$ such that $s \in S$ and $t \in T$
- If f is a flow, then the **net flow** $f(S, T)$ across the cut (S, T) is defined to be:

$$f(S, T) = \sum_{u \in S} \sum_{v \in T} f(u, v) - \sum_{u \in S} \sum_{v \in T} f(v, u) .$$

- The **capacity** of the cut is


$$c(S, T) = \sum_{u \in S} \sum_{v \in T} c(u, v) .$$

- **Example:**

For the shown graph
 $f(S, T) = 19$ and $c(S, T) = 26$

- **Notes:**

- **capacity** counts only the capacities on edges going from S to T .
- **net flow** considers flows in edges on both directions, and it is the same for all cuts.





Minimum Cut of a Flow Network

- A ***minimum cut*** of a network is a cut whose ***capacity*** is minimum over all cuts of the network.
- Let f be a flow in a flow network G with source s and sink t , then *net flow* across any cut is $|f|$.
- *Corollary*: The value of any flow f in a flow network G is bounded from above by the capacity of any cut of G .
- The value of a *maximum flow* in a network equals the capacity of a *minimum cut* of the network.
- The *max-flow min-cut* theorem tells us that a flow is maximum if and only if its residual network contains no augmenting path.



Ford-Fulkerson Algorithm

- The *Ford-Fulkerson* method repeatedly augments the flow along augmenting paths until it has found a maximum flow.


```

FORD-FULKERSON( $G, s, t$ )
1  for each edge  $(u, v) \in G.E$ 
2       $(u, v).f = 0$ 
3  while there exists a path  $p$  from  $s$  to  $t$  in the residual network  $G_f$ 
4       $c_f(p) = \min \{c_f(u, v) : (u, v) \text{ is in } p\}$ 
5      for each edge  $(u, v)$  in  $p$ 
6          if  $(u, v) \in E$ 
7               $(u, v).f = (u, v).f + c_f(p)$ 
8          else  $(v, u).f = (v, u).f - c_f(p)$ 
  
```


- Lines 6–8 update the flow by adding flow when the residual edge is an edge in the original network and subtracting it otherwise.
- When no augmenting paths exist, the flow f is a maximum flow.




Ford-Fulkerson Example (1 of 3)




Input G_f with augmenting path p shaded and $c_f = 4$



Updated G_f with augmenting path p shaded and $c_f = 4$






Updated G_f with augmenting path p shaded and $c_f = 4$



Updated G_f with augmenting path p shaded and $c_f = 7$




Example (2 of 3)




(Repeated)

Updated G_f with augmenting path p shaded and $c_f = 4$



Updated G_f with no augmenting paths exist



Updated G with the maximum flow $f = 23$



Example (3 of 3)

- As you see the **minimum cut** of the given flow network G (which is a cut whose capacity is minimum over all cuts of the network) has **capacity** of 23, which matches the maximum flow found by the *Ford-Fulkerson* algorithm.

