


EECE7205: Fundamentals of Computer Engineering



String Matching




The String-Matching Problem

- Text-editing programs frequently need to find all occurrences of a pattern in the text or to apply automatic spelling correction.
- DNA sequences are searched for particular patterns.
- Internet search engines also need algorithms to find Web pages relevant to queries.
- Efficient algorithms for this problem are called “***string matching***”.
- Given a text represented as an array $T[1..n]$ of length n , the string-matching algorithm is to find in T the occurrences of a pattern $P[1..m]$ of length $m \leq n$.
- Assumption: the elements of P and T are characters drawn from a finite alphabet Σ .
- The character arrays P and T are often called ***strings*** of characters.



String Matching Definitions

- Referring to the figure below, we say that pattern P **occurs with shift** s in text T (or, equivalently, that pattern P **occurs beginning at position** $s + 1$ in text T)
- This means if $0 \leq s \leq n - m$ then $T[s + 1..s + m] = P[1..m]$ (that is, $T[s + j] = P[j]$), for $1 \leq j \leq m$).
- If P occurs with shift s in T , then we call s a **valid shift**; otherwise, we call s an **invalid shift**.
- The **string-matching problem** is the problem of finding **all** valid shifts with which a given pattern P occurs in a given text T .






Notation and Terminology (1 of 2)

- Σ^* = set of all finite-length strings formed using characters from alphabet Σ
- Empty string: ε (also belongs to Σ^*)
- $|x|$ = length of string x
- The **concatenation** of two strings x and y , denoted xy , has length $|x| + |y|$ and consists of the characters from x followed by the characters from y .
- w is a **prefix** of x : $w \sqsubseteq x$ (e.g., $ab \sqsubseteq abcca$)
- w is a **suffix** of x : $w \sqsupseteq x$ (e.g., $cca \sqsupseteq abcca$)
 - In both cases $|w| \leq |x|$
- The prefix and suffix relations are *transitive*.
- The empty string ε is both a suffix and a prefix of any string.



Notation and Terminology (2 of 2)






- We denote the k -character prefix $R[1..k]$ of the string $R[1..n]$ by R_k . Thus, $R_0 = \varepsilon$ and $R_n = R$.
- Using this notation, we can state the string-matching problem as that of:
finding all shifts s in the range $0 \leq s \leq n - m$ such that pattern $P \sqsubseteq T_{s+m}$. Where T is a string of length n and pattern P is of length $m \leq n$.





Overlapping-Suffix Lemma

- Suppose that x , y , and z are strings such that $x \sqsubseteq z$ and $y \sqsubseteq z$.
 If $|x| \leq |y|$, then $x \sqsubseteq y$. If $|x| \geq |y|$, then $y \sqsubseteq x$.
 If $|x| = |y|$, then $x = y$.





The Naive String-Matching Algorithm

- The naive algorithm finds all valid shifts using a loop that checks the condition $P[1..m] = T[s+1..s+m]$ for each of the $n - m + 1$ possible values of s .


NAIVE-STRING-MATCHER(T, P)

```

1   $n = T.length$ 
2   $m = P.length$ 
3  for  $s = 0$  to  $n - m$ 
4      if  $P[1..m] == T[s+1..s+m]$ 
5          print "Pattern occurs with shift"  $s$ 

```

What is the best and worst big O of this algorithm?





The Naive Algorithm Worst Case

- The following example shows the worst-case scenario to search for the pattern *BBC* ($m = 3$) in a text of n characters.

1. BBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBC

BBC 3 comparisons

2. BBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBC

BBC 3 comparisons

3. BBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBC

BBC 3 comparisons

4.

- Total number of comparisons = $m(n-m+1)$
- Time efficiency = $O(nm)$




- Total number of comparisons = $n-m+1$
- Time efficiency = $O(n)$



Finite Automata

- String matching using finite automata avoids testing useless shifts as in the naive pattern-matching algorithm.
- A **finite automaton** M is a 5-tuple $(Q, q_0, A, \Sigma, \delta)$, where
 - Q is a finite set of **states**,
 - $q_0 \in Q$ is the **start state**,
 - $A \subseteq Q$ is a distinguished set of **accepting states**,
 - Σ is a finite **input alphabet**,
 - δ is a function from $Q \times \Sigma$ into Q , called the **transition function** of M .






Finite Automata Example

- In the shown example:

- $Q = \{0, 1\}$,
- $q_0 = 0$,
- $A = \{1\}$
- $\Sigma = \{a, b\}$,
- $\delta(0, a) = 1$
 $\delta(0, b) = 0$
 $\delta(1, a) = 0$
 $\delta(1, b) = 0$

state	input	
	a	b
0	1	0
1	0	0






The Final-State Function

- The function Φ defined on a finite automaton M is called the ***final-state function***.
- For a string $w \in \Sigma^*$, $\Phi(w)$ returns the state in which M ends up after M scans the string w .
- M accepts a string w if and only if $\Phi(w) \in A$ (the set of accepting states).
- We define the function Φ recursively, using the transition function:

$$\phi(\varepsilon) = q_0, \text{ Where } \varepsilon \text{ is the empty string}$$


$$\phi(wa) = \delta(\phi(w), a) \text{ for } w \in \Sigma^*, a \in \Sigma.$$





The Suffix Function

- For a given pattern $P[1 .. m]$, the ***suffix function*** σ maps Σ^* to $\{0, \dots, m\}$ such that $\sigma(x)$ is the length of the longest prefix of P that is also a suffix of x .




- Example:
 - If pattern $P = ab$, we have $\sigma(\epsilon) = 0$,
 - $\sigma(ccaca) = 1$,
 - $\sigma(ccab) = 2$.
 - For a pattern P of length m , we have $\sigma(x) = m$ if and only if $P \sqsubseteq x$.



String-Matching Automata (1 of 2)

- For a given pattern P , we construct a string-matching automaton in a preprocessing step before using it to search the text string.
- The figure illustrates how we construct the automaton for the pattern $P = ababaca$
- Some edges corresponding to failing matches are omitted; by convention, if a state i has no outgoing edge labelled x then $\delta(i, x) = 0$.



- Sample operation:

What is the big O
of the matching
process?

i	—	1	2	3	4	5	6	7	8	9	10	11
$T[i]$	—	a	b	a	b	a	b	a	c	a	b	a
state $\phi(T_i)$	0	1	2	3	4	5	4	5	6	7	2	3



String-Matching Automata (2 of 2)

- We define the string-matching automaton that corresponds to a given pattern $P[1..m]$ as follows:

- The state set Q is $\{0, 1, \dots, m\}$. The start state q_0 is state 0, and state m is the only accepting state.


- The transition function δ is defined by equation

$$\delta(q, x) = \sigma(P_q x) \text{ for any state } q \text{ and character } x$$



- Example1: for $P = \text{ababaca}$, $q = 5$, $x = b \rightarrow \delta(5, b) = \sigma(\text{ababab}) = 4$
- Example2: for $P = \text{ababaca}$, $q = 5$, $x = a \rightarrow \delta(5, a) = \sigma(\text{ababaa}) = 1$
- Example3: for $P = \text{ababaca}$, $q = 5$, $x = c \rightarrow \delta(5, c) = \sigma(\text{ababac}) = 6$

What is the big
O of building
the automata?





Finite Automata Matcher Algorithm

- The following algorithm uses an automaton (represented by its transition function δ) to find occurrences of a pattern P of length m in an input text $T[1..n]$.


FINITE-AUTOMATON-MATCHER(T, δ, m)

```

1   $n = T.length$ 
2   $q = 0$ 
3  for  $i = 1$  to  $n$ 
4       $q = \delta(q, T[i])$ 
5      if  $q == m$ 
6          print "Pattern occurs with shift"  $i - m$ 
  
```

Representing the automaton

- Preprocessing time = $O(m|\Sigma|)$
- Matching time = $\Theta(n)$






The Knuth-Morris-Pratt Algorithm

- This algorithm uses an auxiliary function π that allows us to compute the transition function δ efficiently “on the fly” as needed.
- The array π is pre-computed from the pattern in time $\Theta(m)$ and is stored in an array $\pi[1..m]$.
- For any state $q = 0, 1, \dots, m$ and any character $a \in \Sigma$, the value $\pi[q]$ contains the information we need to compute $\delta(q, a)$.
- Since the array π has only m entries, whereas δ has $\Theta(m |\Sigma|)$ entries, we save a factor of $|\Sigma|$ in the pre-processing time by computing π rather than δ .
 - The idea here comes from the fact that in computing δ , there is no need to consider the characters in Σ that are not in P as they will always bring us back to state 0.

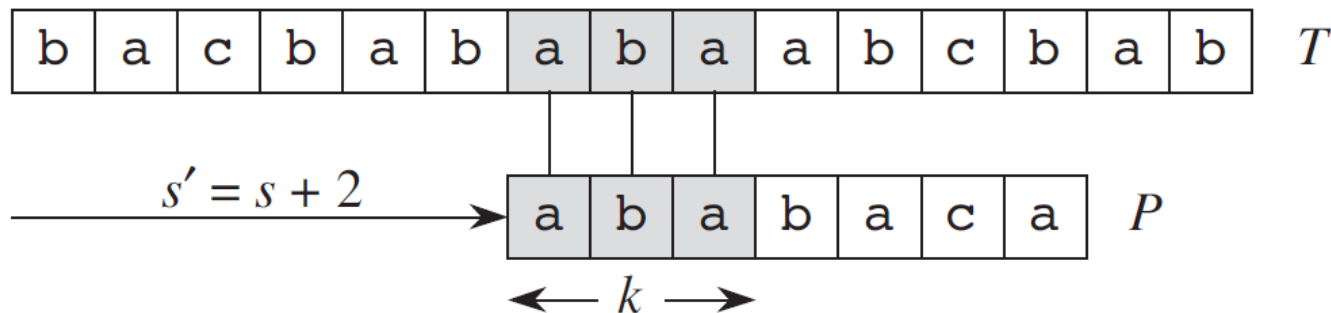


Comparing P With Itself

- The entries in array π are precomputed by comparing the pattern P with itself.
- The shown pattern $P = ababaca$ aligns with a text T so that the first $q = 5$ characters match.



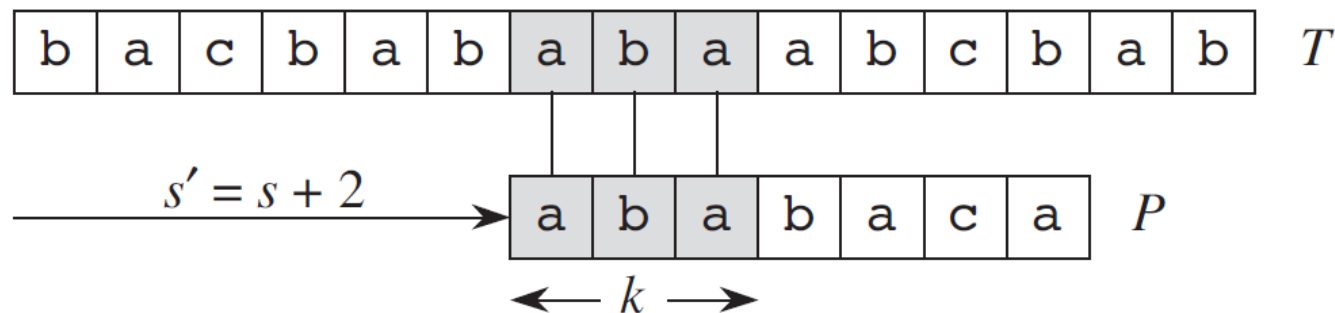

- Using only our knowledge of the 5 matched characters, we can deduce that a shift of $s + 1$ is invalid, but that a shift of $s' = s + 2$ is potentially valid.





An Example of a π 's Entry

- Here, we see that P_3 is the longest prefix of P that is also a proper suffix of P_5 .
- We represent this pre-computed information in the array π , so that $\pi[5]=3$.
- Now we can continue our matching process from where we stopped at T and compare to character $P[\pi[5]+1 = 4]$, which means the new shift $s' = s + (5 - \pi[5])$.
 - Comparing to the automata, this also means that at state q , the next state after a match is $q+1$ and the next state after a mismatch is $\pi[q]$. However here we will need to compare the mismatched character again.






Computing π

- Given a pattern P of m characters where q characters have matched successfully at shift s , the next potentially valid shift is at $s' = s + (q - \pi[q])$ where the **prefix function** for the pattern P is the function $\pi: \{1, 2, \dots, m\} \rightarrow \{0, 1, \dots, m-1\}$ such that:

$$\pi[q] = \max \{k : k < q \text{ and } P_k \sqsubset P_q\} .$$

Which is the longest prefix of P that is also a proper suffix of P_q




Prefix Function Example

i	1	2	3	4	5	6	7
$P[i]$	a	b	a	b	a	c	a
$\pi[i]$	0	0	1	2	3	0	1

The longest prefix of P that is also a proper suffix of P_q

$$\pi[q] = \max \{k : k < q \text{ and } P_k \sqsubset P_q\} .$$






KMP Preprocessing Algorithm

COMPUTE-PREFIX-FUNCTION(P)

```

1   $m = P.length$ 
2  let  $\pi[1..m]$  be a new array
3   $\pi[1] = 0$ 
4   $k = 0$ 
5  for  $q = 2$  to  $m$ 
6      while  $k > 0$  and  $P[k + 1] \neq P[q]$ 
7           $k = \pi[k]$ 
8      if  $P[k + 1] == P[q]$ 
9           $k = k + 1$ 
10      $\pi[q] = k$ 
11 return  $\pi$ 

```



i	1	2	3	4	5	6	7
$P[i]$	a	b	a	b	a	c	a
$\pi[i]$	0	0	1	2	3	0	1

Preprocessing time is $\Theta(m)$

KMP Matcher Algorithm

KMP-MATCHER(T, P)

```



1   $n = T.length$ 
2   $m = P.length$ 
3   $\pi = \text{COMPUTE-PREFIX-FUNCTION}(P)$ 
4   $q = 0$ 
5  for  $i = 1$  to  $n$ 
6      while  $q > 0$  and  $P[q + 1] \neq T[i]$ 
7           $q = \pi[q]$ 
8      if  $P[q + 1] == T[i]$ 
9           $q = q + 1$ 
10     if  $q == m$ 
11         print "Pattern occurs with shift"  $i - m$ 
12      $q = \pi[q]$ 

```

// number of characters matched
// scan the text from left to right
// next character does not match
// next character matches
// is all of P matched?
// look for the next match

i	1	2	3	4	5	6	7
$P[i]$	a	b	a	b	a	c	a
$\pi[i]$	0	0	1	2	3	0	1

Matching time is $\Theta(n)$





Algorithms Running Time

- Except for the naive algorithm, each string-matching algorithm we studied performs some preprocessing based on the pattern and then finds all valid shifts; we call this latter phase “matching.”
- The total running time of each algorithm is the sum of the preprocessing and matching times.

Algorithm	Preprocessing time	Matching time
Naive	0	$O(n m)$
Finite automaton	$O(m \Sigma)$	$\Theta(n)$
Knuth-Morris-Pratt	$\Theta(m)$	$\Theta(n)$